

Observability > Agent > Conversations > e5c16b1f-6da8-4ae7-ad83-20e00fe3a...

Q Search...

ctrl+k

+

⌵

🕒

🌟

👤 Invite

🧩

e5c16b1f-6da8-4ae7-ad83-20e00fe3a51e

🔄

🗨 Annotate

💾 Save

+ Promote to test case

📄

🕒 3m 59s

🕒 Started Sep 2, 2026 · 19:17:02

🕒 Updated Sep 2, 2026 · 19:21:01

🗨 62 calls

📄 4.5M tokens

💰 \$1.89 cost

🟢 OK

🔵 gemini-3.5-flash

AI analysis

2h

→ Conversation consumed \$1.89 across 62 generations with massive token inflation: 4.5M total tokens from just 754K input + 22K output, indicating severe context accumulation or repetition.

→ Top 3 generations account for \$0.449 (24% of total cost) with relatively uniform pricing, suggesting consistent high-token outputs rather than single outliers.

→ 239-second duration with 62 generations averages 3.9 seconds per call; investigate if extended conversation history is being re-processed per generation.

TIMELINE

0s

239.2s

FLOW

DEPENDENCIES

GROUP

Agent

Model

Provider

SORT

Time

Duration

Tokens

Estimated cost

Q Filter...

As {}

#1

🔵 gemini-3.5-flash

1.6k

\$0.008

🟢

🔗

transfer\_to\_agent

#2

🔵 gemini-3.5-flash

16.2k

\$0.025

🟢

🔗

list\_datasources

#3

🔵 gemini-3.5-flash

16.9k

\$0.005

🟢

> METRICS

62 calls · peak 0ms · \$1.89 total

CONVERSATION THREAD

STEP 1

🔵 gemini-3.5-flash

1.6k tok

\$0.0076

User 07:17:02 PM

👤

Investigate the current cinema and video-streaming health across Grafana and Confluent Kafka.

Check for active degradation, identify affected streaming regions or services where telemetry supports it, correlate infrastructure metrics/logs with Kafka events, determine the most likely cause, estimate viewer impact where evidence supports it, and recommend operator actions.

Use only data returned by the available Grafana and Confluent tools. Clearly distinguish verified evidence from hypotheses. Do not invent telemetry.

👤 AGENT 07:17:02 PM